

Movie_Budget_and_Financial_Records

November 21, 2025

1 Movie Budgets and Financial Performance

- 1.1 Do higher film budgets lead to more box office revenue? Let's find out if there's a relationship using the movie budgets and financial performance data scraped from [the-numbers.com](#)

2 Import Statements

```
[54]: import pandas as pd  
import matplotlib.pyplot as plt
```

3 Notebook Presentation

```
[55]: pd.options.display.float_format = '{:,.2f}'.format  
  
from pandas.plotting import register_matplotlib_converters  
register_matplotlib_converters()
```

4 Read the Data

```
[56]: data = pd.read_csv('cost_revenue_dirty.csv')
```

5 Explore and Clean the Data

```
[57]: data
```

	Rank	Release_Date	Movie_Title	USD_Production_Budget	\
0	5293	8/2/1915	The Birth of a Nation	\$110,000	
1	5140	5/9/1916	Intolerance	\$385,907	
2	5230	12/24/1916	20,000 Leagues Under the Sea	\$200,000	
3	5299	9/17/1920	Over the Hill to the Poorhouse	\$100,000	
4	5222	1/1/1925	The Big Parade	\$245,000	
...	
5386	2950	10/8/2018	Meg	\$15,000,000	

```

5387    126    12/18/2018           Aquaman      $160,000,000
5388     96    12/31/2020           Singularity   $175,000,000
5389   1119    12/31/2020       Hannibal the Conqueror   $50,000,000
5390   2517    12/31/2020 Story of Bonnie and Clyde, The   $20,000,000

          USD_Worldwide_Gross USD_Domestic_Gross
0            $11,000,000      $10,000,000
1                  $0             $0
2            $8,000,000      $8,000,000
3            $3,000,000      $3,000,000
4            $22,000,000     $11,000,000
...
5386                 $0             $0
5387                 $0             $0
5388                 $0             $0
5389                 $0             $0
5390                 $0             $0

```

[5391 rows x 6 columns]

5.1 Rows and columns

[58]: `data.shape`

[58]: (5391, 6)

5.2 NaN values

[59]: `data.isna().values.any()`

[59]: False

5.3 Duplicate rows

[60]: `data.duplicated().any() #will return True if there are duplicates`

[60]: False

[61]: `data[data.duplicated()] #show the duplicated rows`

[61]: Empty DataFrame
Columns: [Rank, Release_Date, Movie_Title, USD_Production_Budget,
USD_Worldwide_Gross, USD_Domestic_Gross]
Index: []

5.4 Columns Data Types

```
[62]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5391 entries, 0 to 5390
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Rank              5391 non-null    int64  
 1   Release_Date      5391 non-null    object  
 2   Movie_Title       5391 non-null    object  
 3   USD_Production_Budget  5391 non-null    object  
 4   USD_Worldwide_Gross  5391 non-null    object  
 5   USD_Domestic_Gross  5391 non-null    object  
dtypes: int64(1), object(5)
memory usage: 252.8+ KB
```

5.5 Data Type Conversions

5.5.1 Will convert the `USD_Production_Budget`, `USD_Worldwide_Gross`, and `USD_Domestic_Gross` columns to a numeric format by removing \$ signs and ..

Note that *domestic* in this context refers to the United States.

```
[63]: data.USD_Domestic_Gross = data.USD_Domestic_Gross.astype(str).str.replace('$',  
                           '')
```

```
[64]: data.USD_Domestic_Gross = data.USD_Domestic_Gross.astype(str).str.replace(',',  
                           '')
```

```
[65]: data.USD_Domestic_Gross = pd.to_numeric(data.USD_Domestic_Gross)
```

```
[66]: data.USD_Worldwide_Gross = data.USD_Worldwide_Gross.astype(str).str.  
                           replace('$', '')
```

```
[67]: data.USD_Worldwide_Gross = data.USD_Worldwide_Gross.astype(str).str.  
                           replace(',', '')
```

```
[68]: data.USD_Worldwide_Gross = pd.to_numeric(data.USD_Worldwide_Gross)
```

```
[69]: data.USD_Production_Budget = data.USD_Production_Budget.astype(str).str.  
                           replace('$', '')
```

```
[70]: data.USD_Production_Budget = data.USD_Production_Budget.astype(str).str.  
                           replace(',', '')
```

```
[71]: data.USD_Production_Budget = pd.to_numeric(data.USD_Production_Budget)
```

```
[72]: data
```

```
[72]:      Rank Release_Date          Movie_Title \
0      5293    8/2/1915      The Birth of a Nation
1      5140    5/9/1916           Intolerance
2      5230  12/24/1916  20,000 Leagues Under the Sea
3      5299  9/17/1920  Over the Hill to the Poorhouse
4      5222   1/1/1925        The Big Parade
...
5386  2950  10/8/2018             ...
5387  126   12/18/2018            Meg
5388   96   12/31/2020           Aquaman
5389  1119  12/31/2020        Singularity
5389  1119  12/31/2020       Hannibal the Conqueror
5390  2517  12/31/2020  Story of Bonnie and Clyde, The

      USD_Production_Budget  USD_Worldwide_Gross  USD_Domestic_Gross
0                  110000        11000000          10000000
1                  385907                 0              0
2                  200000        8000000          8000000
3                  100000        3000000          3000000
4                  245000        22000000         11000000
...
5386      15000000             ...
5387     160000000            ...
5388     175000000            ...
5389     50000000            ...
5390     20000000            ...

[5391 rows x 6 columns]
```

5.5.2 Seek and Destroy

```
[73]: chars_to_remove = [',', '$']
columns_to_clean = ['USD_Production_Budget',
                     'USD_Worldwide_Gross',
                     'USD_Domestic_Gross']

for col in columns_to_clean:
    for char in chars_to_remove:
        # Replace each character with an empty string
        data[col] = data[col].astype(str).str.replace(char, "")

    # Convert column to a numeric data type
    data[col] = pd.to_numeric(data[col])
```

```
[74]: data
```

```
[74]:      Rank Release_Date          Movie_Title \
0      5293    8/2/1915      The Birth of a Nation
1      5140    5/9/1916        Intolerance
2      5230  12/24/1916  20,000 Leagues Under the Sea
3      5299  9/17/1920  Over the Hill to the Poorhouse
4      5222   1/1/1925       The Big Parade
...
5386   2950  10/8/2018            ...
5387   126   12/18/2018           Meg
5388    96   12/31/2020         Aquaman
5389  1119   12/31/2020       Singularity
5389  1119   12/31/2020     Hannibal the Conqueror
5390  2517   12/31/2020  Story of Bonnie and Clyde, The

          USD_Production_Budget  USD_Worldwide_Gross  USD_Domestic_Gross
0                  110000          11000000          10000000
1                  385907              0                 0
2                  200000          8000000          8000000
3                  100000          3000000          3000000
4                  245000          22000000          11000000
...
5386                  ...
5387          15000000              0                 0
5387          160000000             0                 0
5388          175000000             0                 0
5389          50000000              0                 0
5390          20000000              0                 0

[5391 rows x 6 columns]
```

5.5.3 Will convert the `Release_Date` column to a Pandas Datetime type.

```
[75]: data.Release_Date = pd.to_datetime(data.Release_Date)
```

```
[76]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5391 entries, 0 to 5390
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   Rank             5391 non-null   int64  
 1   Release_Date     5391 non-null   datetime64[ns]
 2   Movie_Title      5391 non-null   object  
 3   USD_Production_Budget  5391 non-null   int64  
 4   USD_Worldwide_Gross  5391 non-null   int64  
 5   USD_Domestic_Gross  5391 non-null   int64  
dtypes: datetime64[ns](1), int64(4), object(1)
memory usage: 252.8+ KB
```

5.6 Descriptive Statistics

5.6.1 What is the average production budget of the films in the data set?

```
[77]: avg_budget = data.USD_Production_Budget.mean()  
avg_budget
```

```
[77]: 31113737.57837136
```

```
[78]: print(f"The average production budget of the films is {round(avg_budget, 2)}")
```

The average production budget of the films is 31113737.58

5.6.2 What is the average worldwide gross revenue of films?

```
[79]: avg_gross = data.USD_Worldwide_Gross.mean()  
avg_gross
```

```
[79]: 88855421.96271564
```

```
[80]: print(f"The average worldwide gross revenue of the films is  
      ↪{round(avg_gross, 2)}")
```

The average worldwide gross revenue of the films is 88855421.96

5.6.3 Are the bottom 25% of films actually profitable or do they lose money?

5.6.4 What are the highest production budget and highest worldwide gross revenue of any film?

5.6.5 How much revenue did the lowest and highest budget films make?

```
[81]: data.describe()
```

```
[81]:          Rank           Release_Date  USD_Production_Budget  \\\n  count    5,391.00                      5391                  5,391.00\n  mean     2,696.00  2003-09-19 15:02:02.203672704            31,113,737.58\n  min       1.00          1915-08-02 00:00:00                  1,100.00\n  25%     1,348.50          1999-12-02 12:00:00             5,000,000.00\n  50%     2,696.00          2006-06-23 00:00:00            17,000,000.00\n  75%     4,043.50          2011-11-23 00:00:00            40,000,000.00\n  max      5,391.00          2020-12-31 00:00:00           425,000,000.00\n  std      1,556.39                           NaN                40,523,796.88\n\n          USD_Worldwide_Gross  USD_Domestic_Gross\n  count            5,391.00            5,391.00\n  mean        88,855,421.96        41,235,519.44\n  min            0.00              0.00\n  25%        3,865,206.00        1,330,901.50\n  50%        27,450,453.00        17,192,205.00
```

```

75%           96,454,455.00      52,343,687.00
max          2,783,918,982.00    936,662,225.00
std           168,457,757.00     66,029,346.27

```

```
[82]: data[data.USD_Production_Budget == 1100]
```

```

[82]:      Rank Release_Date      Movie_Title  USD_Production_Budget \
2427    5391   2005-05-08  My Date With Drew                      1100

      USD_Worldwide_Gross  USD_Domestic_Gross
2427            181041                181041

```

```
[83]: data[data.USD_Production_Budget == 425000000]
```

```

[83]:      Rank Release_Date Movie_Title  USD_Production_Budget \
3529      1    2009-12-18    Avatar                  425000000

      USD_Worldwide_Gross  USD_Domestic_Gross
3529            2783918982                760507625

```

6 Investigating the Zero Revenue Films

```
[84]: zero_domestic = data[data.USD_Domestic_Gross == 0]
zero_domestic
```

```

[84]:      Rank Release_Date      Movie_Title \
1        5140   1916-05-09      Intolerance
6        4630   1927-12-08      Wings
8        4240   1930-01-01      Hell's Angels
17       4814   1936-10-20  Charge of the Light Brigade, The
27       4789   1941-10-28  How Green Was My Valley
...      ...
5386     2950   2018-10-08      Meg
5387     126    2018-12-18      Aquaman
5388     96     2020-12-31      Singularity
5389    1119   2020-12-31      Hannibal the Conqueror
5390    2517   2020-12-31  Story of Bonnie and Clyde, The

      USD_Production_Budget  USD_Worldwide_Gross  USD_Domestic_Gross
1                  385907                 0                  0
6                  2000000                0                  0
8                  4000000                0                  0
17                 1200000                0                  0
27                 1250000                0                  0
...
5386                15000000                0                  0
5387               160000000               0                  0

```

```

5388           175000000          0          0
5389           50000000          0          0
5390           20000000          0          0

```

[512 rows x 6 columns]

```
[85]: zero_domestic.sort_values('USD_Production_Budget', ascending = False)
```

	Rank	Release_Date	Movie_Title \	
5388	96	2020-12-31	Singularity	
5387	126	2018-12-18	Aquaman	
5384	321	2018-09-03	A Wrinkle in Time	
5385	366	2018-10-08	Amusement Park	
5090	556	2015-12-31	Don Gato, el inicio de la pandilla	
...	
4787	5371	2014-12-31	Stories of Our Lives	
3056	5374	2007-12-31	Tin Can Man	
4907	5381	2015-05-19	Family Motocross	
5006	5389	2015-09-29	Signed Sealed Delivered	
5007	5390	2015-09-29	A Plague So Pleasant	
		USD_Production_Budget	USD_Worldwide_Gross	USD_Domestic_Gross
5388		175000000	0	0
5387		160000000	0	0
5384		103000000	0	0
5385		100000000	0	0
5090		80000000	4547660	0
...	
4787		15000	0	0
3056		12000	0	0
4907		10000	0	0
5006		5000	0	0
5007		1400	0	0

[512 rows x 6 columns]

```
[86]: zero_worldwide = data[data.USD_Worldwide_Gross == 0]
zero_worldwide
```

	Rank	Release_Date	Movie_Title \
1	5140	1916-05-09	Intolerance
6	4630	1927-12-08	Wings
8	4240	1930-01-01	Hell's Angels
17	4814	1936-10-20	Charge of the Light Brigade, The
27	4789	1941-10-28	How Green Was My Valley
...
5386	2950	2018-10-08	Meg

5387	126	2018-12-18	Aquaman		
5388	96	2020-12-31	Singularity		
5389	1119	2020-12-31	Hannibal the Conqueror		
5390	2517	2020-12-31	Story of Bonnie and Clyde, The		
		USD_Production_Budget	USD_Worldwide_Gross	USD_Domestic_Gross	
1		385907	0	0	
6		2000000	0	0	
8		4000000	0	0	
17		1200000	0	0	
27		1250000	0	0	
...
5386		15000000	0	0	
5387		160000000	0	0	
5388		175000000	0	0	
5389		50000000	0	0	
5390		20000000	0	0	

[357 rows x 6 columns]

```
[87]: zero_worldwide.sort_values('USD_Production_Budget', ascending = False)
```

[87]:	Rank	Release_Date	Movie_Title	USD_Production_Budget	\
5388	96	2020-12-31	Singularity	175000000	
5387	126	2018-12-18	Aquaman	160000000	
5384	321	2018-09-03	A Wrinkle in Time	103000000	
5385	366	2018-10-08	Amusement Park	100000000	
5058	880	2015-11-12	The Ridiculous 6	60000000	
...
4787	5371	2014-12-31	Stories of Our Lives	15000	
3056	5374	2007-12-31	Tin Can Man	12000	
4907	5381	2015-05-19	Family Motocross	10000	
5006	5389	2015-09-29	Signed Sealed Delivered	5000	
5007	5390	2015-09-29	A Plague So Pleasant	1400	
USD_Worldwide_Gross		USD_Domestic_Gross			
5388		0		0	
5387		0		0	
5384		0		0	
5385		0		0	
5058		0		0	
...
4787		0		0	
3056		0		0	
4907		0		0	
5006		0		0	
5007		0		0	

```
[357 rows x 6 columns]
```

6.0.1 Filtering on Multiple Conditions

```
[88]: international_releases = data.loc[(data.USD_Domestic_Gross == 0) & (data.  
    ↪USD_Worldwide_Gross != 0)]  
international_releases
```

```
[88]:      Rank Release_Date          Movie_Title  USD_Production_Budget  \  
 71     4310   1956-02-16           Carousel        3380000  
 1579   5087   2001-02-11  Everything Put Together        500000  
 1744   3695   2001-12-31            The Hole       7500000  
 2155   4236   2003-12-31           Nothing       4000000  
 2203   2513   2004-03-31         The Touch      20000000  
...     ...      ...             ...           ...  
 5340   1506   2017-04-14  Queen of the Desert        36000000  
 5348   2225   2017-05-05  Chāi dàn zhuānjiā        23000000  
 5360   4832   2017-07-03        Departure        1100000  
 5372   1856   2017-08-25      Ballerina      30000000  
 5374   4237   2017-08-25  Polina danser sa vie        4000000  
  
      USD_Worldwide_Gross  USD_Domestic_Gross  
 71                 3220              0  
 1579                7890              0  
 1744              10834406             0  
 2155                63180              0  
 2203              5918742              0  
...                 ...              ...  
 5340               1480089              0  
 5348              58807172              0  
 5360                27561              0  
 5372              48048527              0  
 5374                36630              0
```

```
[155 rows x 6 columns]
```

7 or ...

```
[89]: international_releases2 = data.query('USD_Domestic_Gross == 0 and  
    ↪USD_Worldwide_Gross != 0')  
international_releases2
```

```
[89]:      Rank Release_Date          Movie_Title  USD_Production_Budget  \  
 71     4310   1956-02-16           Carousel        3380000  
 1579   5087   2001-02-11  Everything Put Together        500000
```

1744	3695	2001-12-31	The Hole	7500000
2155	4236	2003-12-31	Nothing	4000000
2203	2513	2004-03-31	The Touch	20000000
...
5340	1506	2017-04-14	Queen of the Desert	36000000
5348	2225	2017-05-05	Chāi dàn zhuānjīā	23000000
5360	4832	2017-07-03	Departure	1100000
5372	1856	2017-08-25	Ballerina	30000000
5374	4237	2017-08-25	Polina danser sa vie	4000000
			USD_Worldwide_Gross	USD_Domestic_Gross
71		3220		0
1579		7890		0
1744		10834406		0
2155		63180		0
2203		5918742		0
...
5340		1480089		0
5348		58807172		0
5360		27561		0
5372		48048527		0
5374		36630		0

[155 rows x 6 columns]

7.0.1 Unreleased Films

```
[90]: # Date of Data Collection
scrape_date = pd.Timestamp('2018-5-1')
```

```
[91]: future_releases = data[data.Release_Date >= scrape_date]
future_releases
```

	Rank	Release_Date	Movie_Title \
5384	321	2018-09-03	A Wrinkle in Time
5385	366	2018-10-08	Amusement Park
5386	2950	2018-10-08	Meg
5387	126	2018-12-18	Aquaman
5388	96	2020-12-31	Singularity
5389	1119	2020-12-31	Hannibal the Conqueror
5390	2517	2020-12-31	Story of Bonnie and Clyde, The

	USD_Production_Budget	USD_Worldwide_Gross	USD_Domestic_Gross
5384	103000000	0	0
5385	100000000	0	0
5386	15000000	0	0
5387	160000000	0	0

5388	175000000	0	0
5389	50000000	0	0
5390	20000000	0	0

```
[92]: data_clean = data.drop(future_releases.index)
data_clean
```

```
[92]: Rank Release_Date Movie_Title \
0    5293   1915-08-02 The Birth of a Nation
1    5140   1916-05-09 Intolerance
2    5230   1916-12-24 20,000 Leagues Under the Sea
3    5299   1920-09-17 Over the Hill to the Poorhouse
4    5222   1925-01-01 The Big Parade
...
5379 1295   2017-10-02 John Wick: Chapter Two
5380 70     2017-10-03 Kong: Skull Island
5381 94     2017-12-05 King Arthur: Legend of the Sword
5382 1254   2017-12-05 Snatched
5383 2521   2017-12-31 The Thousand Miles

USD_Production_Budget USD_Worldwide_Gross USD_Domestic_Gross
0                  110000          11000000        10000000
1                  385907              0             0
2                  200000          8000000         8000000
3                  100000          3000000         3000000
4                  245000          22000000        11000000
...
5379            40000000      166893990       92029184
5380            185000000     561137727       168052812
5381            175000000     140012608       39175066
5382            42000000      57850343        45850343
5383            20000000              0             0
```

[5384 rows x 6 columns]

7.0.2 Films that Lost Money

```
[93]: losing_money = data_clean.loc[data_clean.USD_Worldwide_Gross < data_clean.
    ↪USD_Production_Budget]
losing_money
```

```
[93]: Rank Release_Date Movie_Title \
1    5140   1916-05-09 Intolerance
6    4630   1927-12-08 Wings
8    4240   1930-01-01 Hell's Angels
15   4738   1936-05-02 Modern Times
17   4814   1936-10-20 Charge of the Light Brigade, The
```

```

...
5371 4901 2017-07-28          An Inconvenient Sequel
5373 2161 2017-08-25          Tulip Fever
5374 4237 2017-08-25          Polina danser sa vie
5381    94 2017-12-05 King Arthur: Legend of the Sword
5383 2521 2017-12-31          The Thousand Miles

      USD_Production_Budget  USD_Worldwide_Gross  USD_Domestic_Gross
1                  385907              0                  0
6                 2000000              0                  0
8                 4000000              0                  0
15                1500000          165049            163245
17                1200000              0                  0
...
5371                ...          130874            130874
5373                ...              0                  0
5374                ...          36630                  0
5381                ...        140012608            39175066
5383                ...          2000000              0

```

[2007 rows x 6 columns]

```
[94]: len(losing_money)/len(data_clean)
```

```
[94]: 0.37277117384843983
```

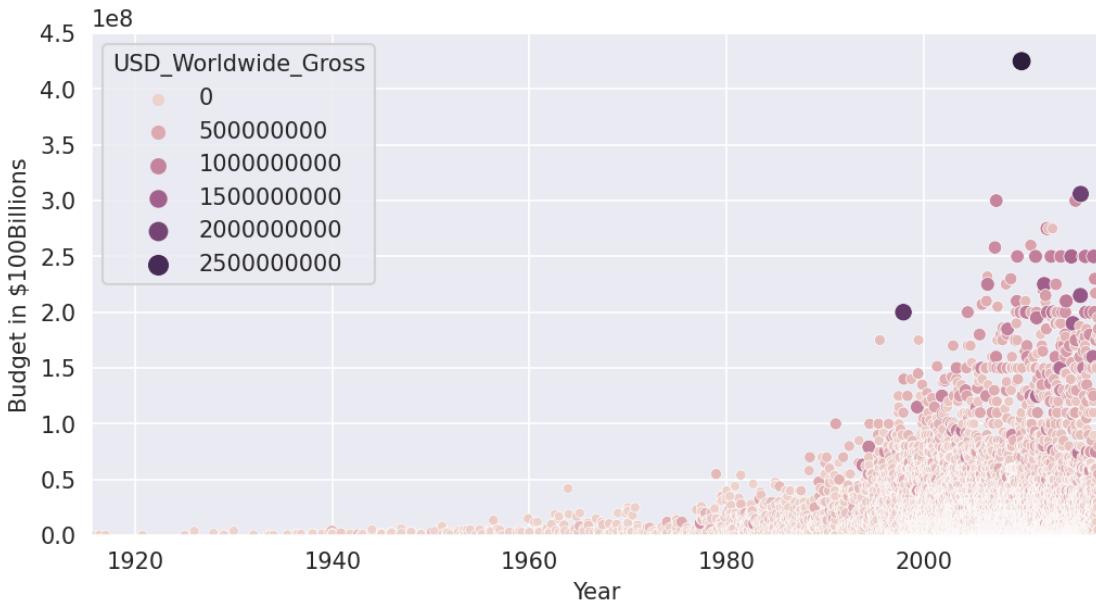
8 Seaborn for Data Viz: Bubble Charts

```
[95]: import seaborn as sns
```

```
[96]: plt.figure(figsize = (8,4), dpi=150)

with sns.axes_style('darkgrid'):
    ax = sns.scatterplot(data=data_clean,
                          x='Release_Date',
                          y='USD_Production_Budget',
                          hue = 'USD_Worldwide_Gross',
                          size = 'USD_Worldwide_Gross')
    ax.set(ylim = (0, 450000000),
           xlim = (data_clean.Release_Date.min(), data_clean.Release_Date.max()),
           ylabel = 'Budget in $100Billions',
           xlabel = 'Year')

plt.show()
```



9 Converting Years to Decades

```
[97]: dt_index = pd.DatetimeIndex(data_clean.Release_Date)
dt_index
```

```
[97]: DatetimeIndex(['1915-08-02', '1916-05-09', '1916-12-24', '1920-09-17',
                   '1925-01-01', '1925-12-30', '1927-12-08', '1929-01-02',
                   '1930-01-01', '1931-12-31',
                   ...
                   '2017-08-25', '2017-09-06', '2017-09-06', '2017-10-02',
                   '2017-10-02', '2017-10-02', '2017-10-03', '2017-12-05',
                   '2017-12-05', '2017-12-31'],
                  dtype='datetime64[ns]', name='Release_Date', length=5384,
                  freq=None)
```

9.0.1 Turn dates into years

```
[98]: years = dt_index.year
years
```

```
[98]: Index([1915, 1916, 1916, 1920, 1925, 1925, 1927, 1929, 1930, 1931,
           ...
           2017, 2017, 2017, 2017, 2017, 2017, 2017, 2017, 2017, 2017],
           dtype='int32', name='Release_Date', length=5384)
```

9.0.2 Turn years into decades with floor division. $1995 / 10 = 199.5$, but $1995 // 10 = 199$. So $199 * 10 = 1990$

```
[99]: decades = (years // 10)*10  
decades
```

```
[99]: Index([1910, 1910, 1910, 1920, 1920, 1920, 1920, 1930, 1930,  
          ...  
          2010, 2010, 2010, 2010, 2010, 2010, 2010, 2010, 2010],  
          dtype='int32', name='Release_Date', length=5384)
```

9.0.3 Add a “Decades” column in the data frame

```
[100]: data_clean['Decade'] = decades  
data_clean
```

```
[100]:      Rank Release_Date           Movie_Title  \n 0      5293  1915-08-02      The Birth of a Nation\n 1      5140  1916-05-09      Intolerance\n 2      5230  1916-12-24  20,000 Leagues Under the Sea\n 3      5299  1920-09-17  Over the Hill to the Poorhouse\n 4      5222  1925-01-01      The Big Parade\n ...  ...  ...  ...  
5379  1295  2017-10-02  John Wick: Chapter Two\n5380    70  2017-10-03      Kong: Skull Island\n5381    94  2017-12-05  King Arthur: Legend of the Sword\n5382  1254  2017-12-05      Snatched\n5383  2521  2017-12-31      The Thousand Miles  
  
      USD_Production_Budget  USD_Worldwide_Gross  USD_Domestic_Gross  Decade\n 0                  110000          11000000          10000000      1910\n 1                  385907              0                0      1910\n 2                  200000          8000000          8000000      1910\n 3                  100000          3000000          3000000      1920\n 4                  245000          22000000          11000000      1920\n...  ...  ...  ...  
5379            40000000          166893990          92029184      2010\n5380          185000000          561137727          168052812      2010\n5381          175000000          140012608          39175066      2010\n5382          42000000          57850343          45850343      2010\n5383          20000000              0                0      2010  
  
[5384 rows x 7 columns]
```

9.0.4 Separate the “old” (before 1969) and “New” (1970s onwards) Films

```
[101]: old_films = data_clean[data_clean.Decade <= 1969]
new_films = data_clean[data_clean.Decade > 1969]
old_films
```

	Rank	Release_Date	Movie_Title		
0	5293	1915-08-02	The Birth of a Nation		
1	5140	1916-05-09	Intolerance		
2	5230	1916-12-24	20,000 Leagues Under the Sea		
3	5299	1920-09-17	Over the Hill to the Poorhouse		
4	5222	1925-01-01	The Big Parade		
..		
148	2375	1969-10-15	Paint Your Wagon		
149	3831	1969-10-24	Butch Cassidy and the Sundance Kid		
150	2175	1969-12-16	Hello, Dolly		
151	3613	1969-12-18	On Her Majesty's Secret Service		
152	4195	1969-12-19	Topaz		
		USD_Production_Budget	USD_Worldwide_Gross	USD_Domestic_Gross	Decade
0		110000	11000000	10000000	1910
1		385907	0	0	1910
2		200000	8000000	8000000	1910
3		100000	3000000	3000000	1920
4		245000	22000000	11000000	1920
..	
148		20000000	31678778	31678778	1960
149		6000000	102308900	102308900	1960
150		24000000	33208099	33208099	1960
151		8000000	82000000	22800000	1960
152		4000000	6000000	6000000	1960

[153 rows x 7 columns]

```
[102]: new_films
```

	Rank	Release_Date	Movie_Title
153	2159	1970-01-01	Waterloo
154	2270	1970-01-01	Darling Lili
155	3136	1970-01-01	Patton
156	3277	1970-01-01	The Molly Maguires
157	4265	1970-01-01	M*A*S*H
..
5379	1295	2017-10-02	John Wick: Chapter Two
5380	70	2017-10-03	Kong: Skull Island
5381	94	2017-12-05	King Arthur: Legend of the Sword
5382	1254	2017-12-05	Snatched
5383	2521	2017-12-31	The Thousand Miles

```

USD_Production_Budget  USD_Worldwide_Gross  USD_Domestic_Gross  Decade
153          25000000          0           0   1970
154          22000000  5000000           5000000   1970
155          12000000  62500000           62500000   1970
156          11000000          2200000           2200000   1970
157          3500000  81600000           81600000   1970
...
5379         40000000  166893990           92029184   2010
5380         185000000  561137727           168052812   2010
5381         175000000  140012608           39175066   2010
5382         42000000  57850343           45850343   2010
5383         20000000          0           0   2010

```

[5231 rows x 7 columns]

[103]: old_films.describe()

	Rank	Release_Date	USD_Production_Budget	\
count	153.00	153	153.00	
mean	4,274.77	1954-06-10 04:04:42.352941184	4,611,297.65	
min	1,253.00	1915-08-02 00:00:00	100,000.00	
25%	3,973.00	1946-01-01 00:00:00	1,250,000.00	
50%	4,434.00	1956-12-23 00:00:00	2,900,000.00	
75%	4,785.00	1964-10-22 00:00:00	5,000,000.00	
max	5,299.00	1969-12-19 00:00:00	42,000,000.00	
std	742.14	NaN	5,713,648.85	

	USD_Worldwide_Gross	USD_Domestic_Gross	Decade
count	153.00	153.00	153.00
mean	30,419,634.38	22,389,473.87	1,949.15
min	0.00	0.00	1,910.00
25%	5,273,000.00	5,000,000.00	1,940.00
50%	10,000,000.00	10,000,000.00	1,950.00
75%	33,208,099.00	28,350,000.00	1,960.00
max	390,525,192.00	198,680,470.00	1,960.00
std	54,931,828.93	32,641,752.41	12.72

[104]: old_films.sort_values('USD_Production_Budget', ascending = False).head(10)

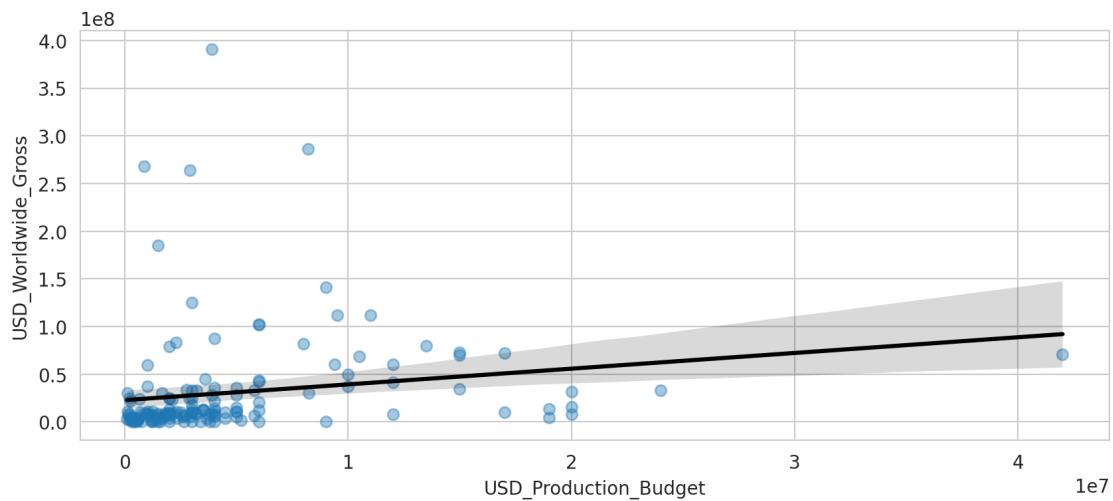
	Rank	Release_Date	Movie_Title	USD_Production_Budget	\
109	1253	1963-12-06	Cleopatra	42000000	
150	2175	1969-12-16	Hello, Dolly	24000000	
143	2465	1969-01-01	Sweet Charity	20000000	
118	2425	1965-02-15	The Greatest Story Ever Told	20000000	
148	2375	1969-10-15	Paint Your Wagon	20000000	
110	2552	1964-01-01	The Fall of the Roman Empire	19000000	

98	2546	1962-08-11	Mutiny on The Bounty	19000000
114	2670	1964-10-22	My Fair Lady	17000000
102	2698	1963-01-01	55 Days at Peking	17000000
125	2831	1966-10-10	Hawaii	15000000
109		USD_Worldwide_Gross	USD_Domestic_Gross	Decade
109		71000000	57000000	1960
150		33208099	33208099	1960
143		8000000	8000000	1960
118		15473333	15473333	1960
148		31678778	31678778	1960
110		4750000	4750000	1960
98		13680000	13680000	1960
114		72070955	72000000	1960
102		10000000	10000000	1960
125		34562222	34562222	1960

10 Seaborn Regression Plots

```
[105]: plt.figure(figsize = (10,4), dpi=200)
```

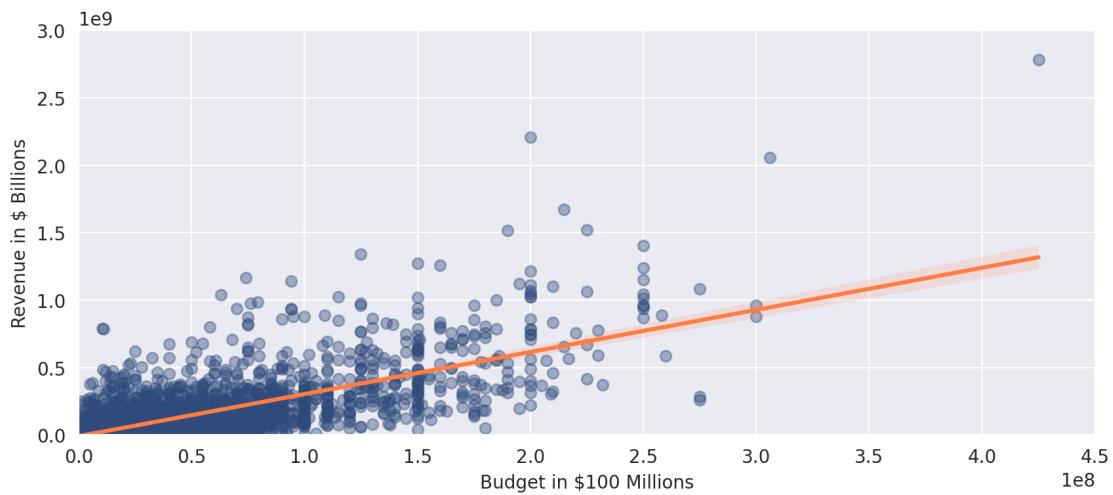
```
with sns.axes_style('whitegrid'):
    ax = sns.regplot(data = old_films,
                      x = 'USD_Production_Budget',
                      y = 'USD_Worldwide_Gross',
                      scatter_kws = {'alpha' : 0.4},
                      line_kws = {'color' : 'black'})
```



```
[106]: plt.figure(figsize = (10,4), dpi=200)

with sns.axes_style('darkgrid'):
    ax = sns.regplot(data = new_films,
                      x = 'USD_Production_Budget',
                      y = 'USD_Worldwide_Gross',
                      color = '#2f4b7c',
                      scatter_kws = {'alpha' : 0.4},
                      line_kws = {'color' : '#ff7c43'})

    ax.set(ylim = (0, 3000000000),
           xlim = (0, 450000000),
           ylabel = 'Revenue in $ Billions',
           xlabel = 'Budget in $100 Millions')
```



11 Running Regression with scikit-learn

$$\hat{REVENUE} = \theta_0 + \theta_1 BUDGET$$

```
[107]: from sklearn.linear_model import LinearRegression
```

11.0.1 Will run a linear regression for the `old_films` and Calculate the intercept, slope and r-squared.

11.0.2 How much of the variance in movie revenue does the linear model explain in this case?

```
[108]: regression = LinearRegression()
```

11.0.3 Explanatory variable (Feature in ML)

```
[109]: X = pd.DataFrame(new_films, columns = ['USD_Production_Budget'])
```

11.0.4 Response variable (Target in ML)

```
[110]: y = pd.DataFrame(new_films, columns = ['USD_Worldwide_Gross'])
```

11.0.5 Creating DataFrames because LinearRegression doesn't like to receive Pandas Series

12 Will find the best fit line

```
[111]: regression.fit(X, y)
```

```
[111]: LinearRegression()
```

```
[163]: regression.intercept_ #Theta zero
```

```
[163]: array([-8650768.00661042])
```

```
[164]: regression.coef_ #Theta one
```

```
[164]: array([[3.12259592]])
```

```
[168]: regression.score(X, y) #R^2
```

```
[168]: 0.5577032617720403
```

Our model explains about 56% of the variance in movie revenue.

12.0.1 Regression

```
[113]: X = pd.DataFrame(old_films, columns = ['USD_Production_Budget'])
y = pd.DataFrame(old_films, columns = ['USD_Worldwide_Gross'])
```

```
[114]: regression.fit(X, y)
```

```
[114]: LinearRegression()
```

```
[115]: regression.intercept_
```

```
[115]: array([22821538.63508039])  
[116]: regression.coef_  
[116]: array([[1.64771314]])  
[117]: regression.score(X, y)  
[117]: 0.02937258620576877  
[118]: print(f'The intercept is : {regression.intercept_}')  
print(f'The slope is : {regression.coef_}')  
print(f'The R2 is : {regression.score(X, y)}')  
print(f'That means our model explains {round(regression.score(X, y) * 100)}% of  
the variance in movie revenue.')
```

The intercept is : [22821538.63508039]
The slope is : [[1.64771314]]
The R² is : 0.02937258620576877
That means our model explains 3% of the variance in movie revenue.

12.0.2 How much global revenue does our model estimate for a film with a budget of i.e. \$350 million?

```
[119]: budget = 350000000  
[120]: revenue_estimate = regression.intercept_[0] + regression.coef_[0,0] * budget  
[121]: revenue_estimate  
[121]: 599521139.0388364  
[122]: revenue_estimate = round(revenue_estimate, -6) # me to -6 can round stin 6h  
# taxi aristera tis ypodia stolis. Sto ekatommio dld, 10^6  
[123]: revenue_estimate  
[123]: 600000000.0  
[124]: print(f'The revenue estimate for a movie with a budget of ${budget} is expected  
to be around ${revenue_estimate:.10}')
```

The revenue estimate for a movie with a budget of \$350000000 is expected to be around \$600000000.0